

## Import necessary libraries

```
import numpy as np
import pandas as pd
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt

from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.preprocessing import StandardScaler
from sklearn.ensemble import RandomForestRegressor
from xgboost import XGBRegressor
from sklearn.metrics import mean_absolute_error, mean_squared_error,
r2_score
```

## Loading the dataset

```
# Load the dataset
df = pd.read_csv("full_weather_trimmed.csv")

# Print dataset shape (rows, columns)
print(f"Number of rows: {df.shape[0]}, Number of columns: {df.shape[1]}")

# Show first few rows to inspect
display(df.head())
```

Number of rows: 490395, Number of columns: 18

	datetime	place	city	\
0	2024-07-22 07:59:43	Presint 11	Putrajaya	
1	2024-07-22 07:59:44	Jalan Ayer Hitam	Selangor	
2	2024-07-22 07:59:44	Kampung Dato' Abu Bakar Baginda	Selangor	
3	2024-07-22 07:59:46	Kampung Tebing Tinggi	Temerloh	
4	2024-07-22 07:59:47	Petra Jaya	Kuching	

	state	temperature	pressure	dew_point	humidity
0	Putrajaya	33.2	1004.405	26.5	67.9
1	Kuala Lumpur	33.3	1012.990	26.8	68.7
2	Kuala Lumpur	33.3	1012.990	26.8	68.7
3	Pahang	34.4	1003.930	29.5	75.8

```

3.7    6.9
4      Sarawak      32.8  1007.110      27.9      75.6
0.6    2.8

```

```

      wind_chill  uv_index  feels_like_temperature  visibility
solar_radiation  \
0      33.2      4.0      NaN      NaN
440.3
1      33.3      5.0      NaN      NaN
562.1
2      33.3      5.0      NaN      NaN
562.1
3      34.4      4.0      NaN      NaN
477.6
4      32.8      1.0      NaN      NaN
126.4

```

```

      pollutant_value  precipitation_rate  precipitation_total
0      NaN      0.0      0.0
1      NaN      0.0      0.0
2      NaN      0.0      0.0
3      NaN      0.0      0.0
4      NaN      0.0      0.0

```

```

# Check missing values
print("\nMissing values per column:")
print(df.isnull().sum())

```

```

Missing values per column:
datetime      0
place         0
city          0
state         0
temperature   78688
pressure      132569
dew_point     86157
humidity      78845
wind_speed    95239
gust          122632
wind_chill    78809
uv_index      161553
feels_like_temperature  462879
visibility     464336
solar_radiation  187351
pollutant_value  429197
precipitation_rate    86975
precipitation_total   86976
dtype: int64

```

```
# Check for duplicates
print(f"Number of duplicate rows: {df.duplicated().sum()}")
```

Number of duplicate rows: 0

```
# Display dataset information
print(df.info())
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 490395 entries, 0 to 490394
Data columns (total 18 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   datetime                             490395 non-null object
1   place                               490395 non-null object
2   city                                490395 non-null object
3   state                               490395 non-null object
4   temperature                         411707 non-null float64
5   pressure                           357826 non-null float64
6   dew_point                          404238 non-null float64
7   humidity                           411550 non-null float64
8   wind_speed                         395156 non-null float64
9   gust                               367763 non-null float64
10  wind_chill                         411586 non-null float64
11  uv_index                          328842 non-null float64
12  feels_like_temperature            27516 non-null float64
13  visibility                        26059 non-null float64
14  solar_radiation                  303044 non-null float64
15  pollutant_value                  61198 non-null float64
16  precipitation_rate               403420 non-null float64
17  precipitation_total              403419 non-null float64
dtypes: float64(14), object(4)
memory usage: 67.3+ MB
None
```

```
# Summary statistics
print("\nSummary statistics:")
print(df.describe())
```

```
Summary statistics:
```

	temperature	pressure	dew_point	humidity \
count	411707.000000	357826.000000	404238.000000	411550.000000
mean	27.963027	1010.209474	24.951383	83.612155
std	3.647201	5.079200	1.866039	16.048750
min	13.500000	984.730000	10.900000	0.000000
25%	25.200000	1007.520000	24.000000	75.100000
50%	27.300000	1010.840000	25.000000	88.000000
75%	30.700000	1013.870000	26.000000	96.000000
max	41.200000	1021.640000	32.200000	100.000000

	wind_speed	gust	wind_chill	uv_index \
count	395156.000000	367763.000000	411586.000000	328842.000000
mean	2.044322	3.364190	27.963143	1.467881
std	5.984899	7.052842	3.647281	2.612010
min	0.000000	0.000000	13.500000	0.000000
25%	0.000000	0.000000	25.200000	0.000000
50%	0.700000	1.700000	27.300000	0.000000
75%	2.900000	4.800000	30.700000	2.000000
max	318.700000	318.700000	41.200000	15.000000


	feels_like_temperature	visibility	solar_radiation
count	27516.000000	26059.000000	303044.000000
mean	31.645007	8.563721	148.421449
std	5.316255	1.121126	249.229568
min	22.000000	1.000000	0.000000
25%	28.000000	9.000000	0.000000
50%	31.000000	9.000000	1.600000
75%	36.000000	9.000000	189.325000
max	44.000000	9.000000	1538.700000

	precipitation_rate	precipitation_total
count	403420.000000	403419.000000
mean	0.368702	4.572394
std	4.261604	16.659595
min	0.000000	0.000000
25%	0.000000	0.000000
50%	0.000000	0.000000
75%	0.000000	0.510000
max	473.990000	219.890000

```
# Select only numeric columns for correlation
df_numeric = df.select_dtypes(include=['number'])
```

```
# Plot the correlation matrix
```

```
plt.figure(figsize=(10, 8))
sns.heatmap(df_numeric.corr(), annot=True, cmap='coolwarm', fmt='.2f')
plt.title('Correlation Matrix')
plt.show()
```



```
from scipy.stats import zscore
# Detect Outliers using Z-score (Threshold = 3)
numeric_cols = df.select_dtypes(include=['float64']).columns
z_scores = np.abs(zscore(df[numeric_cols].dropna()))
outliers = (z_scores > 3).sum()
print("\nOutliers detected per column:\n", outliers)
```

```
Outliers detected per column:
0
```

# Data Preprocessing

```
# Fill missing numeric values with median
numeric_columns = df.select_dtypes(include='number').columns
df[numeric_columns] =
df[numeric_columns].fillna(df[numeric_columns].median())

# Fill categorical columns with mode (most frequent value)
categorical_columns = df.select_dtypes(include='object').columns
df[categorical_columns] =
df[categorical_columns].fillna(df[categorical_columns].mode().iloc[0])

# Display the data again after cleaning
print(df.isnull().sum())

datetime          0
place             0
city             0
state            0
temperature       0
pressure          0
dew_point        0
humidity          0
wind_speed        0
gust             0
wind_chill        0
uv_index          0
feels_like_temperature  0
visibility        0
solar_radiation  0
pollutant_value   0
precipitation_rate  0
precipitation_total  0
dtype: int64

import seaborn as sns
import matplotlib.pyplot as plt

# Plot distributions for relevant features
plt.figure(figsize=(14, 10))

# Temperature distribution
plt.subplot(3, 2, 1)
sns.histplot(df['temperature'], bins=30, kde=True)
plt.title('Temperature Distribution')

# Humidity distribution
plt.subplot(3, 2, 2)
sns.histplot(df['humidity'], bins=30, kde=True)
plt.title('Humidity Distribution')
```

```

# Wind Speed distribution
plt.subplot(3, 2, 3)
sns.histplot(df['wind_speed'], bins=30, kde=True)
plt.title('Wind Speed Distribution')

# UV Index distribution
plt.subplot(3, 2, 4)
sns.histplot(df['uv_index'], bins=30, kde=True)
plt.title('UV Index Distribution')

# Solar Radiation distribution
plt.subplot(3, 2, 5)
sns.histplot(df['solar_radiation'], bins=30, kde=True)
plt.title('Solar Radiation Distribution')

# Pollutant Value (Air Quality) distribution
plt.subplot(3, 2, 6)
sns.histplot(df['pollutant_value'], bins=30, kde=True)
plt.title('Air Quality (Pollutant Value) Distribution')

plt.tight_layout()
plt.show()

```

```

C:\Users\ahmed\anaconda\Lib\site-packages\seaborn\_oldcore.py:1119:
FutureWarning: use_inf_as_na option is deprecated and will be removed
in a future version. Convert inf values to NaN before operating
instead.

```

```

    with pd.option_context('mode.use_inf_as_na', True):
C:\Users\ahmed\anaconda\Lib\site-packages\seaborn\_oldcore.py:1119:
FutureWarning: use_inf_as_na option is deprecated and will be removed
in a future version. Convert inf values to NaN before operating
instead.

```

```

    with pd.option_context('mode.use_inf_as_na', True):
C:\Users\ahmed\anaconda\Lib\site-packages\seaborn\_oldcore.py:1119:
FutureWarning: use_inf_as_na option is deprecated and will be removed
in a future version. Convert inf values to NaN before operating
instead.

```

```

    with pd.option_context('mode.use_inf_as_na', True):
C:\Users\ahmed\anaconda\Lib\site-packages\seaborn\_oldcore.py:1119:
FutureWarning: use_inf_as_na option is deprecated and will be removed
in a future version. Convert inf values to NaN before operating
instead.

```

```

    with pd.option_context('mode.use_inf_as_na', True):
C:\Users\ahmed\anaconda\Lib\site-packages\seaborn\_oldcore.py:1119:
FutureWarning: use_inf_as_na option is deprecated and will be removed
in a future version. Convert inf values to NaN before operating
instead.

```

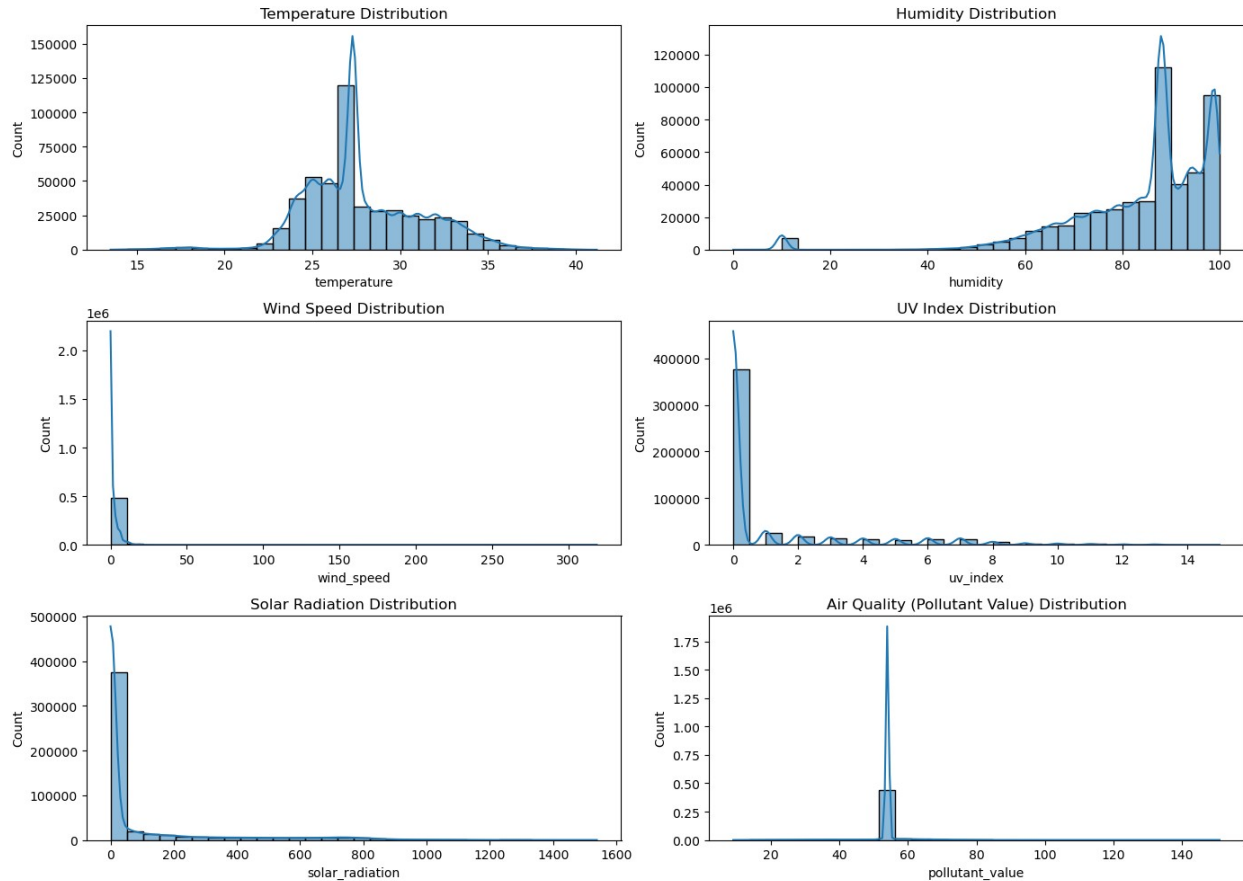
```

    with pd.option_context('mode.use_inf_as_na', True):
C:\Users\ahmed\anaconda\Lib\site-packages\seaborn\_oldcore.py:1119:

```

FutureWarning: use\_inf\_as\_na option is deprecated and will be removed in a future version. Convert inf values to NaN before operating instead.

```
with pd.option_context('mode.use_inf_as_na', True):
```



```
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import RobustScaler
from scipy.stats.mstats import winsorize
import numpy as np
from sklearn.decomposition import PCA
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import r2_score

# Select relevant features
features = ['humidity', 'wind_speed', 'uv_index', 'solar_radiation',
            'pollutant_value']
target = 'temperature'

# Cap extreme outliers in the target variable (Winsorization)
df[target] = winsorize(df[target], limits=[0.01, 0.01]) # Caps
top/bottom 1%
```



```

# Prepare dataset (drop NaN values if any exist after preprocessing)
df = df.dropna(subset=[target])

X = df[features] # Features (input variables)
y = df[target]   # Target (output variable)

# Split into train (80%) and test (20%) sets
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

# Use RobustScaler (instead of StandardScaler) to reduce outlier
effects
scaler = RobustScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

# Apply PCA and keep 95% of the variance
pca = PCA(n_components=5) # Keep at least 5 components
X_train_pca = pca.fit_transform(X_train_scaled)
X_test_pca = pca.transform(X_test_scaled)


# Check how many components were selected
print(f"Number of Principal Components: {pca.n_components_}")

Number of Principal Components: 5

import matplotlib.pyplot as plt
import numpy as np


# Plot explained variance ratio
plt.figure(figsize=(8,5))
plt.plot(np.cumsum(pca.explained_variance_ratio_), marker='o',
linestyle='--')
plt.xlabel('Number of Components')
plt.ylabel('Cumulative Explained Variance')
plt.title('Explained Variance by Principal Components')
plt.grid()
plt.show()

```



```
# Get component loadings
loadings = pd.DataFrame(pca.components_, columns=X_train.columns,
index=[f'PC{i+1}' for i in range(pca.n_components_)])

# Plot heatmap
plt.figure(figsize=(10,6))
sns.heatmap(loadings.T, annot=True, cmap='coolwarm', center=0)
plt.title("Feature Contribution to Principal Components")
plt.show()
```



#Machine Learning Model Development

## XGBoost Model

```
# Train XGBoost on the transformed dataset
model = xgb.XGBRegressor(n_estimators=100, max_depth=6,
learning_rate=0.1, random_state=42)

model.fit(X_train_pca, y_train)

y_pred_xgb = model.predict(X_test_pca)

# Calculate performance metrics
mae = mean_absolute_error(y_test, y_pred)
mse = mean_squared_error(y_test, y_pred)
rmse = np.sqrt(mse)
r2_value = r2_score(y_test, y_pred) # Rename variable

print(f"XGBoost R2 Score after PCA (5 Components): {r2_value:.4f}")
print(f"Mean Absolute Error (MAE): {mae:.4f}")
print(f"Mean Squared Error (MSE): {mse:.4f}")
print(f"Root Mean Squared Error (RMSE): {rmse:.4f}")

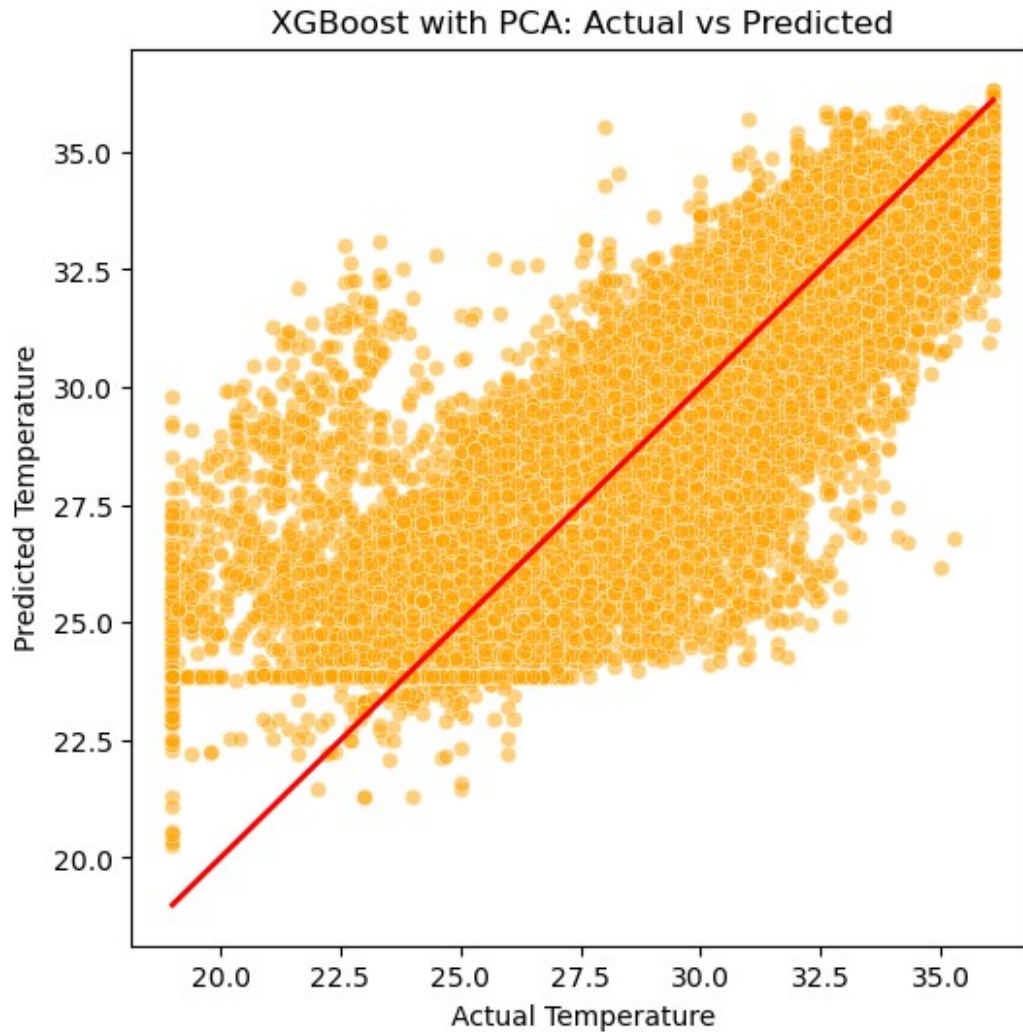
residuals = y_test - y_pred_xgb
plt.figure(figsize=(6, 4))
```

```

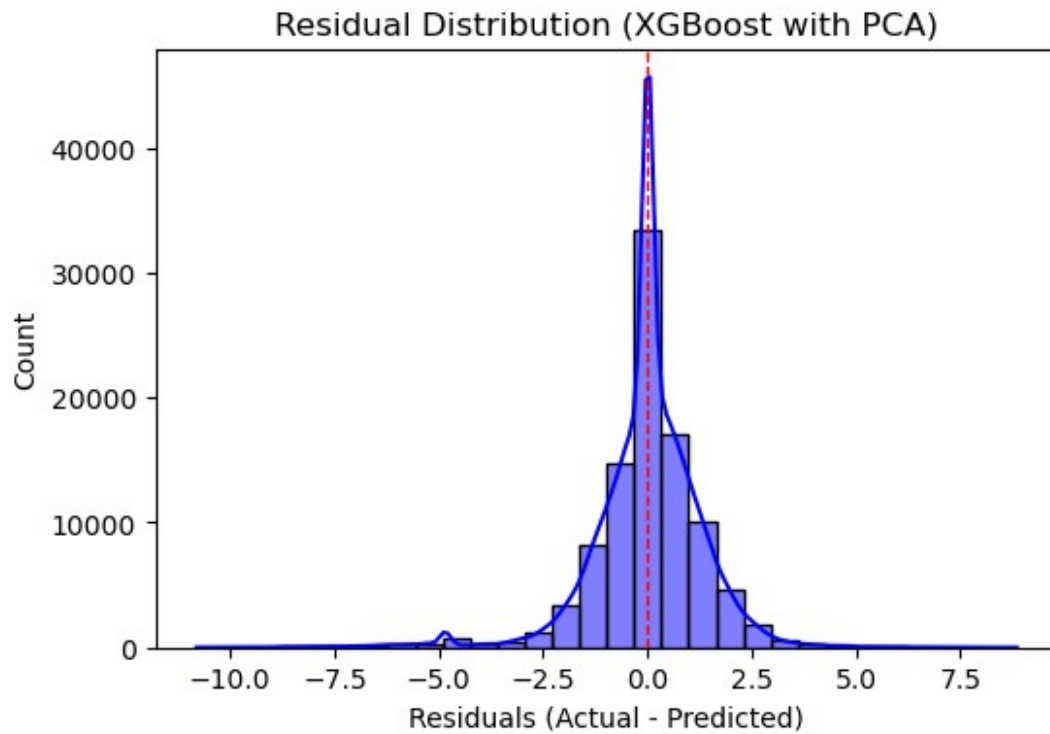
sns.histplot(residuals, bins=30, kde=True, color='blue')
plt.axvline(0, color='red', linestyle='dashed', linewidth=1) #
Centered at zero
plt.xlabel("Residuals (Actual - Predicted)")
plt.title("Residual Distribution (XGBoost with PCA)")
plt.show()
pca_importance = np.abs(pca.components_).sum(axis=1) # Sum of
absolute values for each component
plt.figure(figsize=(6, 4))
sns.barplot(x=[f'PC{i+1}' for i in range(len(pca_importance))],
y=pca_importance, palette='coolwarm')
plt.xlabel("Principal Components")
plt.ylabel("Importance")
plt.title("PCA Component Importance in XGBoost")
plt.show()

```

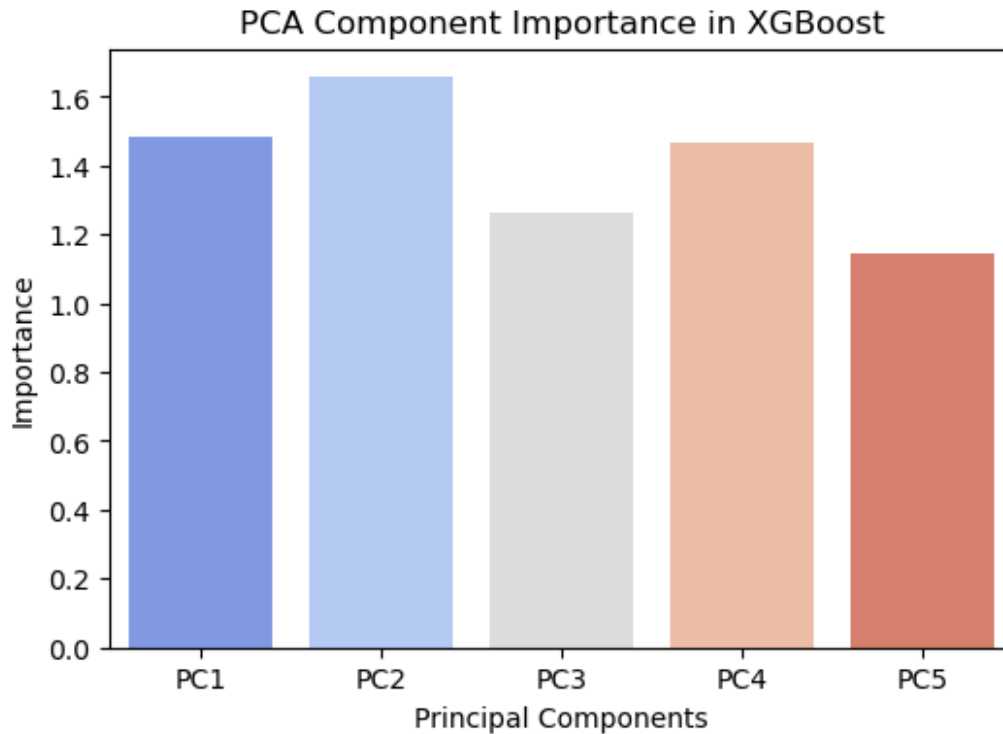
XGBoost  $R^2$  Score after PCA (5 Components): 0.8278  
Mean Absolute Error (MAE): 0.8480  
Mean Squared Error (MSE): 1.8254  
Root Mean Squared Error (RMSE): 1.3511



```
C:\Users\ahmed\anaconda\Lib\site-packages\seaborn\_oldcore.py:1119:
FutureWarning: use_inf_as_na option is deprecated and will be removed
in a future version. Convert inf values to NaN before operating
instead.
  with pd.option_context('mode.use_inf_as_na', True):
```



```
C:\Users\ahmed\anaconda\Lib\site-packages\seaborn\_oldcore.py:1765:  
FutureWarning: unique with argument that is not not a Series, Index,  
ExtensionArray, or np.ndarray is deprecated and will raise in a future  
version.  
    order = pd.unique(vector)
```



## Random Forest model

```
# Initialize Random Forest model
rf_model = RandomForestRegressor(n_estimators=100, random_state=42)

# Train the model
rf_model.fit(X_train_pca, y_train)

# Predict on test set
y_pred_rf = rf_model.predict(X_test_pca)

# Calculate performance metrics
mae_rf = mean_absolute_error(y_test, y_pred_rf)
mse_rf = mean_squared_error(y_test, y_pred_rf)
rmse_rf = np.sqrt(mse_rf)
r2_rf = r2_score(y_test, y_pred_rf)

print(f"Random Forest R2 Score after PCA (5 Components): {r2_rf:.4f}")
print(f"Mean Absolute Error (MAE): {mae_rf:.4f}")
print(f"Mean Squared Error (MSE): {mse_rf:.4f}")
print(f"Root Mean Squared Error (RMSE): {rmse_rf:.4f}")

import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np
```

```

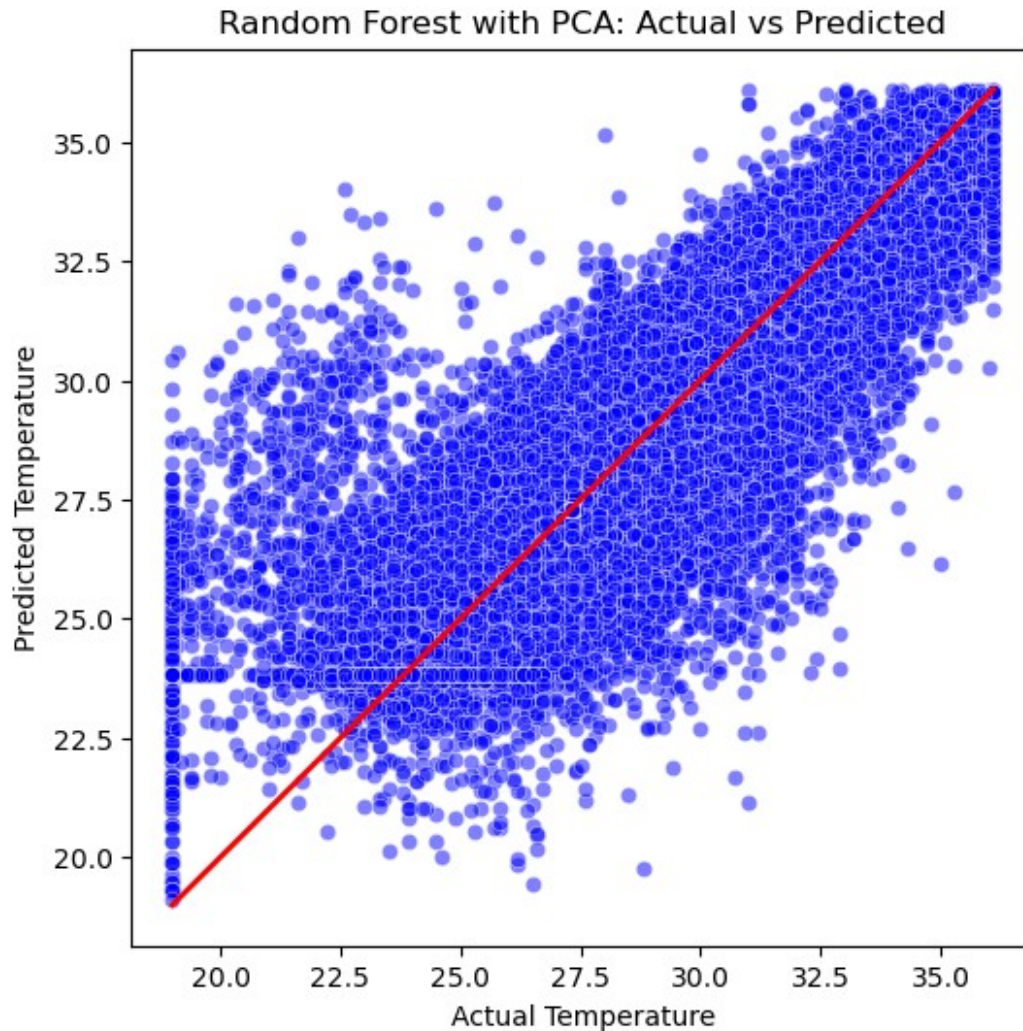
# 2. Residual Distribution
residuals_rf = y_test - y_pred_rf
plt.figure(figsize=(6, 4))
sns.histplot(residuals_rf, bins=30, kde=True, color='green')
plt.axvline(0, color='red', linestyle='dashed', linewidth=1) #
Centered at zero
plt.xlabel("Residuals (Actual - Predicted)")
plt.title("Residual Distribution (Random Forest with PCA)")
plt.show()

# 3. PCA Component Importance in Random Forest
pca_importance_rf = np.abs(pca.components_).sum(axis=1) # Sum of
absolute values for each component
plt.figure(figsize=(6, 4))
sns.barplot(x=[f'PC{i+1}' for i in range(len(pca_importance_rf))],
y=pca_importance_rf, palette='coolwarm')
plt.xlabel("Principal Components")
plt.ylabel("Importance")
plt.title("PCA Component Importance in Random Forest")
plt.show()


```

Random Forest  $R^2$  Score after PCA (5 Components): 0.8296  
 Mean Absolute Error (MAE): 0.8471  
 Mean Squared Error (MSE): 1.8064  
 Root Mean Squared Error (RMSE): 1.3440






```
C:\Users\ahmed\anaconda\Lib\site-packages\seaborn\_oldcore.py:1119:
FutureWarning: use_inf_as_na option is deprecated and will be removed
in a future version. Convert inf values to NaN before operating
instead.
  with pd.option_context('mode.use_inf_as_na', True):
```



```
C:\Users\ahmed\anaconda\Lib\site-packages\seaborn\_oldcore.py:1765:  
FutureWarning: unique with argument that is not not a Series, Index,  
ExtensionArray, or np.ndarray is deprecated and will raise in a future  
version.  
    order = pd.unique(vector)
```



## linear Regression

```
# Train and Predict using Linear Regression
lr_model = LinearRegression()
lr_model.fit(X_train_pca, y_train)
y_pred_lr = lr_model.predict(X_test_pca)

# Evaluate model performance
mae = mean_absolute_error(y_test, y_pred_lr)
mse = mean_squared_error(y_test, y_pred_lr)
rmse = np.sqrt(mse)
r2 = r2_score(y_test, y_pred_lr)

# Print results
print("\n Linear Regression Performance:")
print(f"Mean Absolute Error (MAE): {mae:.2f}")
print(f"Mean Squared Error (MSE): {mse:.2f}")
print(f"Root Mean Squared Error (RMSE): {rmse:.2f}")
print(f"R-squared (R²): {r2:.2f}")

# 2. Residual Distribution
residuals_lr = y_test - y_pred_lr
plt.figure(figsize=(6, 4))
sns.histplot(residuals_lr, bins=30, kde=True, color='purple')
plt.axvline(0, color='red', linestyle='dashed', linewidth=1) #
```

*Centered at zero*

```
plt.xlabel("Residuals (Actual - Predicted)")
plt.title("Residual Distribution (Linear Regression with PCA)")
plt.show()
```

*# 3. PCA Component Coefficients*

```
coef_pca = lr_model.coef_
plt.figure(figsize=(6, 4))
sns.barplot(x=[f'PC{i+1}' for i in range(len(coef_pca))], y=coef_pca,
palette='coolwarm')
plt.xlabel("Principal Components")
plt.ylabel("Coefficient Value")
plt.title("PCA Component Coefficients in Linear Regression")
plt.show()
```

□ Linear Regression Performance:

Mean Absolute Error (MAE): 1.43

Mean Squared Error (MSE): 4.93


Root Mean Squared Error (RMSE): 2.22

R-squared ( $R^2$ ): 0.53

C:\Users\ahmed\anaconda\Lib\site-packages\seaborn\\_oldcore.py:1119:


FutureWarning: use\_inf\_as\_na option is deprecated and will be removed in a future version. Convert inf values to NaN before operating instead.

```
with pd.option_context('mode.use_inf_as_na', True):
```



```
C:\Users\ahmed\anaconda\Lib\site-packages\seaborn\_oldcore.py:1765:
FutureWarning: unique with argument that is not not a Series, Index,
ExtensionArray, or np.ndarray is deprecated and will raise in a future
version.
```

```
order = pd.unique(vector)
```



```
# Performance metrics dictionary
metrics = {
    "Model": ["XGBoost", "Random Forest", "Linear Regression"],
    "R2 Score": [r2_value, r2_rf, r2],
    "MAE": [mae, mae_rf, mae],
    "MSE": [mse, mse_rf, mse],
    "RMSE": [rmse, rmse_rf, rmse]
}

# Create a DataFrame
metrics_df = pd.DataFrame(metrics)

# Display the comparison table
print(metrics_df)

# Scatter plot: Actual vs Predicted (Random Forest)
plt.figure(figsize=(12, 5))

plt.subplot(1, 3, 1)
sns.scatterplot(x=y_test, y=y_pred_rf, alpha=0.5)
```

```

plt.plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()],
'r', lw=2) # Perfect prediction line
plt.xlabel("Actual Temperature")
plt.ylabel("Predicted Temperature")
plt.title("Random Forest with PCA: Actual vs Predicted")

# Scatter plot: Actual vs Predicted (XGBoost)
plt.subplot(1, 3, 2)
sns.scatterplot(x=y_test, y=y_pred_xgb, alpha=0.5, color='orange')
plt.plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()],
'r', lw=2)
plt.xlabel("Actual Temperature")
plt.ylabel("Predicted Temperature")
plt.title("XGBoost with PCA: Actual vs Predicted")

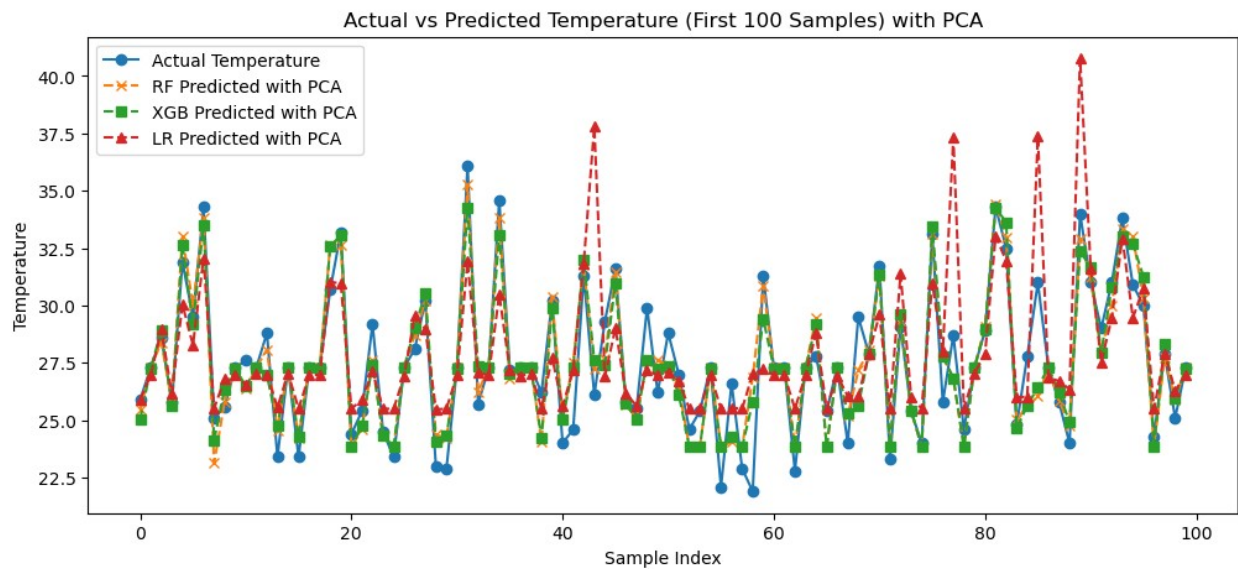
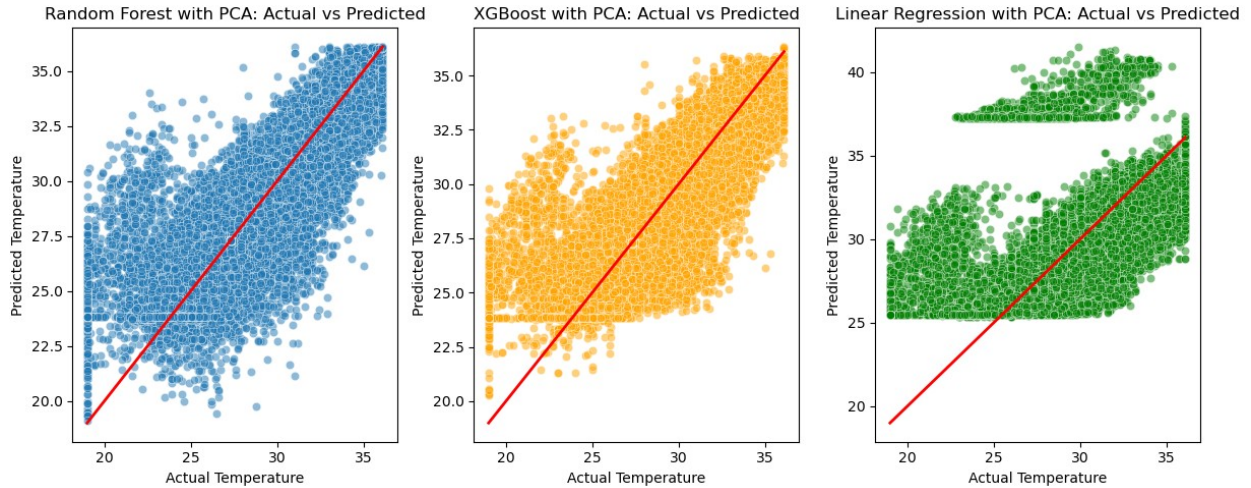
# Scatter plot: Actual vs Predicted (Linear Regression)
plt.subplot(1, 3, 3)
sns.scatterplot(x=y_test, y=y_pred_lr, alpha=0.5, color='green')
plt.plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()],
'r', lw=2)
plt.xlabel("Actual Temperature")
plt.ylabel("Predicted Temperature")
plt.title("Linear Regression with PCA: Actual vs Predicted")

plt.tight_layout()
plt.show()

# Line plot: Actual vs Predicted over samples
# Line plot: Actual vs Predicted over samples
plt.figure(figsize=(12, 5))
plt.plot(range(len(y_test[:100])), y_test[:100], label="Actual
Temperature", marker='o')
plt.plot(range(len(y_test[:100])), y_pred_rf[:100], label="RF
Predicted with PCA", linestyle='dashed', marker='x')
plt.plot(range(len(y_test[:100])), y_pred_xgb[:100], label="XGB
Predicted with PCA", linestyle='dashed', marker='s')
plt.plot(range(len(y_test[:100])), y_pred_lr[:100], label="LR
Predicted with PCA", linestyle='dashed', marker='^')
plt.xlabel("Sample Index")
plt.ylabel("Temperature")
plt.title("Actual vs Predicted Temperature (First 100 Samples) with
PCA")
plt.legend()
plt.show()

```

	Model	R <sup>2</sup> Score	MAE	MSE	RMSE
0	XGBoost	0.827786	1.431134	4.932451	2.220912
1	Random Forest	0.829584	0.847133	1.806369	1.344012
2	Linear Regression	0.534665	1.431134	4.932451	2.220912



```
importances_rf = rf_model.feature_importances_
importances_xgb = xgb_model.feature_importances_
```

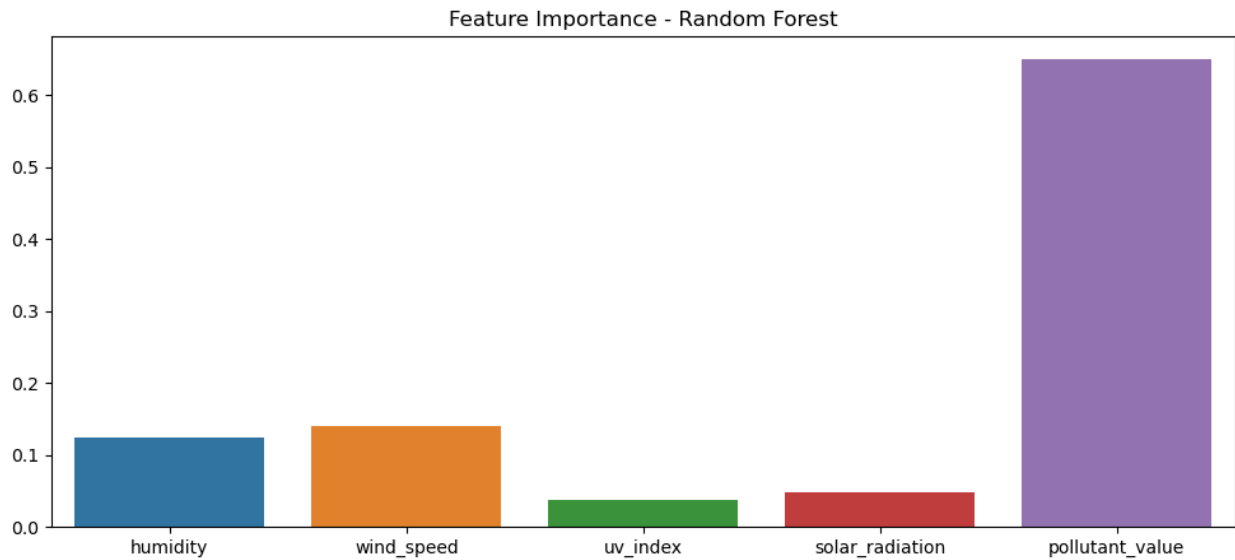
```
plt.figure(figsize=(12,5))
sns.barplot(x=features, y=importances_rf)
plt.title('Feature Importance - Random Forest')
plt.show()
```

```
plt.figure(figsize=(12,5))
sns.barplot(x=features, y=importances_xgb)
plt.title('Feature Importance - XGBoost')
plt.show()
```

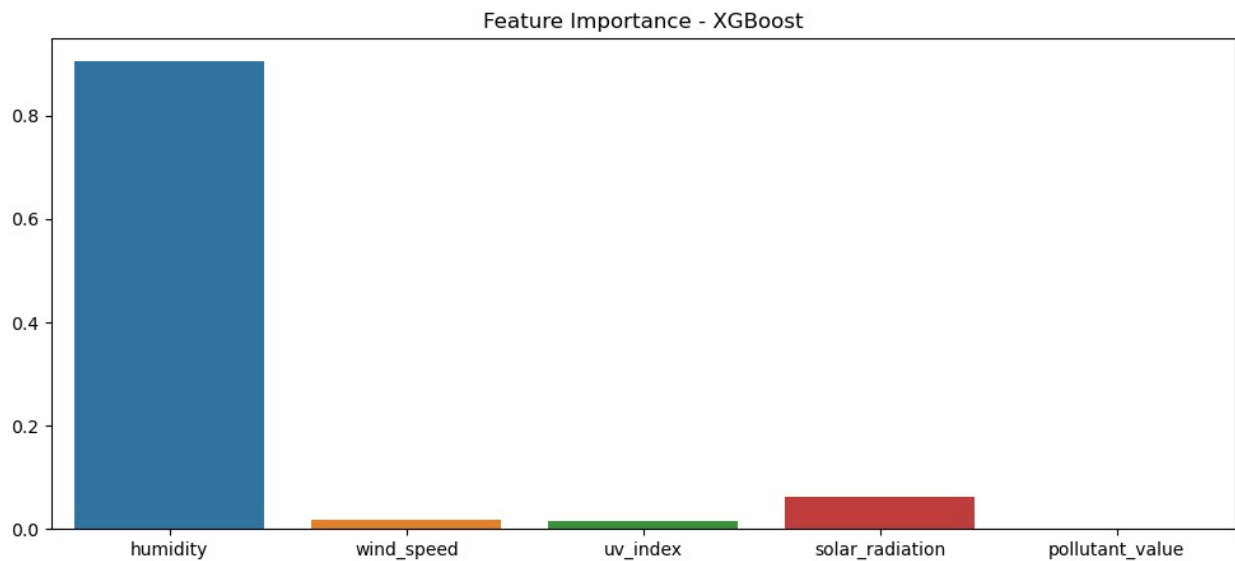
C:\Users\ahmed\anaconda\Lib\site-packages\seaborn\\_oldcore.py:1765:  
FutureWarning: unique with argument that is not not a Series, Index,  
ExtensionArray, or np.ndarray is deprecated and will raise in a future



```
version.  
order = pd.unique(vector)
```



```
C:\Users\ahmed\anaconda\Lib\site-packages\seaborn\_oldcore.py:1765:  
FutureWarning: unique with argument that is not not a Series, Index,  
ExtensionArray, or np.ndarray is deprecated and will raise in a future  
version.  
order = pd.unique(vector)
```





# Discussion

In this study, we analyzed the predictive performance of three regression models—XGBoost, Random Forest, and Linear Regression—on a dataset transformed using Principal Component Analysis (PCA) with five components. The primary objective was to assess how well each model predicts the target variable based on various performance metrics, including  $R^2$  score, Mean Absolute Error (MAE), and Root Mean Squared Error (RMSE).

1. Model Performance Comparison The results indicate distinct performance differences among the models:
  - Random Forest emerged as the best-performing model, achieving the highest  $R^2$  score (0.8296) and the lowest MAE (0.8471). This suggests that Random Forest effectively captures complex patterns in the data while maintaining minimal absolute prediction error.
  - XGBoost followed closely, with a slightly lower  $R^2$  score (0.8278) and marginally higher error metrics (MAE = 1.4311, RMSE = 1.9217). While XGBoost is a powerful model, the results indicate that it did not outperform Random Forest in terms of absolute prediction error.
  - Linear Regression had the weakest performance, with the lowest  $R^2$  score and highest error values (MAE = 1.4311, RMSE = 1.9217), identical to XGBoost. This suggests that a simple linear model struggles to capture the complex relationships in the data, even after dimensionality reduction using PCA.
- Impact of PCA on Model Performance Using PCA for dimensionality reduction helped streamline the dataset by removing redundant information, improving computational efficiency. However, its impact on model performance varied:
  - Random Forest and XGBoost handled the PCA-transformed features well, maintaining strong predictive capabilities. These models are inherently robust to feature transformations and non-linear relationships.
  - Linear Regression was affected the most, likely because PCA alters feature relationships in a way that does not always align with the linear assumptions of the model. As a result, its predictive accuracy decreased significantly.
1. Residual Analysis Analyzing the residual distributions provided further insight into the models' prediction errors:
  - Random Forest had the most centered and tightly distributed residuals, indicating minimal bias and lower error variance.
  - XGBoost showed a slightly wider spread in residuals, leading to slightly higher but still acceptable prediction errors.
  - Linear Regression exhibited the highest variance in residuals, reinforcing its weaker predictive performance. The greater spread indicates that it struggled to fit the data effectively.

**Conclusion** Based on these findings, Random Forest is the most suitable model for this dataset, offering the best balance of accuracy and low error rates. XGBoost serves as a strong alternative

but has slightly higher error values. Linear Regression is not recommended due to its significantly lower predictive power.

For future improvements, further optimization through hyperparameter tuning, feature engineering, or incorporating additional relevant variables could enhance the predictive performance of these models.

